




# Wordcount + PageRank + Apache Spark + GKE

## Kubernetes Project



PageRank

Shoumya Singh

# Table of Content

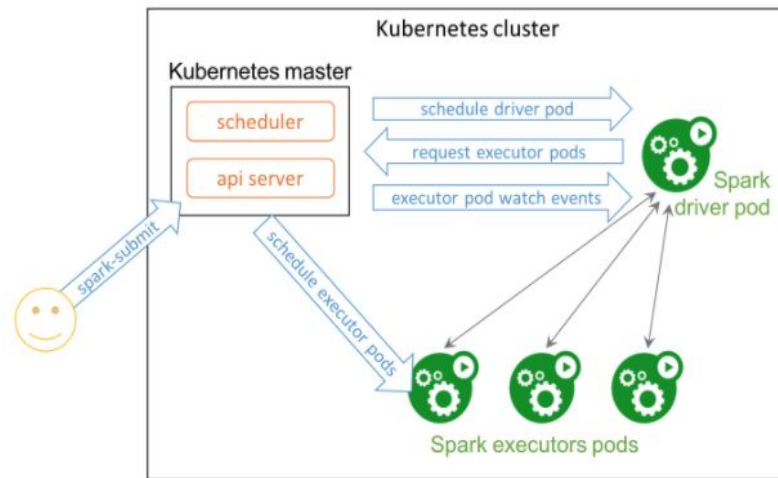
- ❖ **Introduction**
- ❖ **Design**
  - Project Description
  - Concepts - Basic Ideas
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**



kubernetes

# Introduction

- ❑ **Kubernetes** is a portable, extensible, open-source platform for managing **containerized workloads and services**, that facilitates both declarative configuration and automation.
- ❑ Kubernetes provides you with a **framework** to run **distributed systems** resiliently. It takes care of scaling and failover for your application, provides deployment patterns, and more.
- ❑ **Apache Spark** with native Kubernetes support combines the best of the two prominent **open source** projects – Apache Spark, a **framework** for large-scale data processing; and Kubernetes.
- ❑ **PySpark** is an **interface** for Apache Spark in **Python**. It not only allows you to write Spark applications using Python APIs, but also provides the PySpark shell for interactively analyzing your data in a **distributed environment**.



# Table of Content

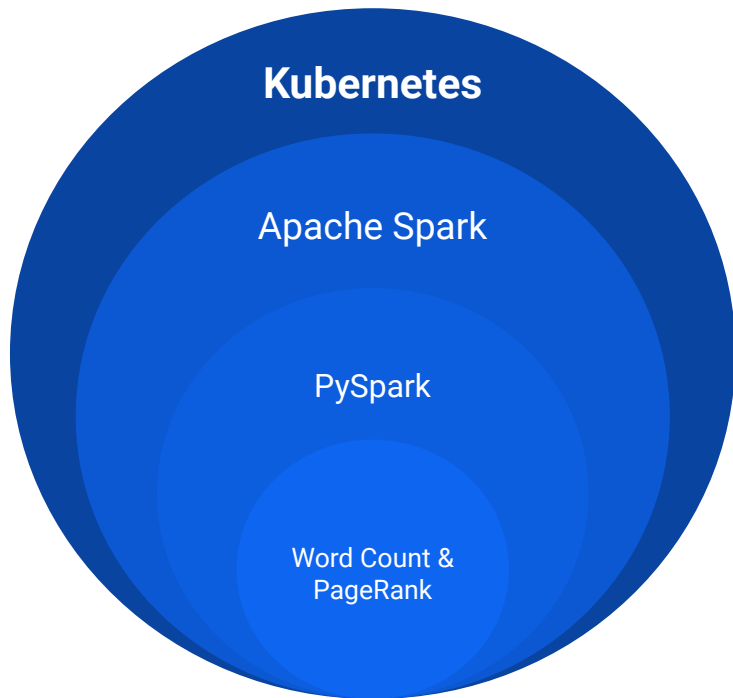
- ❖ **Introduction**
- ❖ **Design**
  - **Project Description**
  - **Concepts - Basic Ideas**
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**

# Project Description



kubernetes

→ Using PySpark to implement Word Count and PageRank on Apache Spark running on Kubernetes.



→ **Technologies Used**

- ◆ Google Kubernetes Engine
- ◆ Apache Spark - PySpark

→ **Applications**

- ◆ Wordcount
- ◆ PageRank



kubernetes

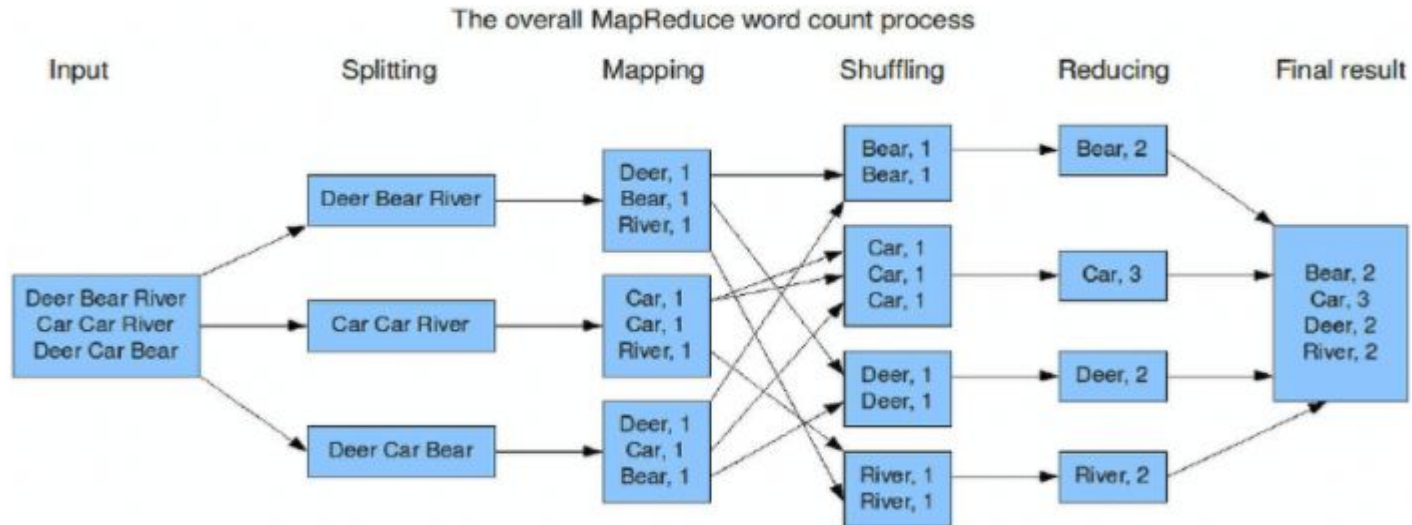
# Table of Content

- ❖ **Introduction**
- ❖ **Design**
  - Project Description
  - Concepts - Basic Ideas
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**



# Wordcount

- ❑ **Wordcount** is a [simple spark application](#) where it counts how often each word appears in a collection of text documents.
- ❑ A distributed computing framework that can run WordCount efficiently in parallel at scale can likely handle much larger and more interesting compute problems.
- ❑ **MapReduce** is a processing technique and a program model for distributed computing based on java. The **MapReduce** algorithm contains two important tasks, namely Map and Reduce. Map takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key/value pairs).




# PageRank



kubernetes

- ❑ **PageRank** is one of many **algorithms** Google uses to work out which order to display search results.
- ❑ The co-founders of Google, Sergey Brin and Larry Page developed the PageRank algorithm in 1996 at Stanford University.
- ❑ Increasing the PageRank score of a web page will mean that page is displayed higher than other pages in a search engine listing, which means more visitors and therefore potentially more customers or money generated from a web page.
- ❑ The formula used to calculate PageRank is:  
$$PR(A) = (1 - d) + d ( PR(t1) / C(t1) + \dots + PR(tn)/C(tn) )$$
- ❑ PageRank is an iterative algorithm which means you repeat the calculation for each page multiple times until the values eventually settle on the final PageRank scores for each page.






# PageRank Calculation Example



kubernetes



→ The **initial PageRank value** for each webpage is 1.

- $PR(A) = 1$
- $PR(B) = 1$
- $PR(C) = 1$

→ Page **B** has a link to pages **C** and **A**

→ Page **C** has a link to page **A**

→ Page **D** has links to all **three pages**

1. **A's PageRank** is  
 $PR(A) = (1-d) + d * (PR(B) / 2 + PR(C) / 1 + PR(D) / 3)$
2. **B's PageRank** is  
 $PR(B) = (1-d) + d * (PR(D) / 3)$
3. **C's PageRank** is  
 $PR(C) = (1-d) + d * (PR(B) / 2 + PR(D) / 3)$
4. **C's PageRank** is  
 $PR(D) = 1-d$
5. **Damping factor** is 0.85 (default value)

# PageRank Calculation Example



kubernetes

Then after 1st iteration

## A. Output

- Page B would transfer half of its existing value, or 0.5, to page A and the other half, or 0.5, to page C.
- Page C would transfer all of its existing value, 1, to the only page it links to, A.
- Since D had three outbound links, it would transfer one third of its existing value, or approximately 0.33, to A.

## C. Observation of PageRank

- The more inputs a node has the more its PageRank will increase in the long run.
- A node does not have input will have
  - constant PageRank:  $1-d$
  - the smallest PageRank

## B. Input

- $PR(A)$   
 $= (1-d) + d * (PR(B) / 2 + PR(C) / 1 + PR(D) / 3)$   
 $= (1-0.85) + 0.85 * (0.5 + 1 + 0.33)$   
 $= 1.71$
- $PR(B)$   
 $= (1-d) + d * (PR(D) / 3)$   
 $= (1-0.85) + 0.85 * 0.33$   
 $= 0.43$
- $PR(C)$   
 $= (1-d) + d * (PR(B) / 2 + PR(D) / 3)$   
 $= (1-0.85) + 0.85 * (0.5 + 0.33)$   
 $= 0.86$
- $PR(D)$   
 $= 1-d$   
 $= 0.15$



# Table of Content

- ❖ **Introduction**
- ❖ **Design**
  - Project Description
  - Concepts - Basic Ideas
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**

# Implementation



kubernetes

## 1. Create a cluster on Google Kubernetes Engine with

- `gcloud container clusters create spark --num-nodes=1 --machine-type=e2-highmem-2 --region=us-west1`

```
NAME      LOCATION  MASTER_VERSION  MASTER_IP      MACHINE_TYPE  NODE_VERSION  NUM_NODES  STATUS
spark     us-west1  1.18.16-gke.502  35.185.198.199  e2-highmem-2  1.18.16-gke.502  3          RUNNING
singh19566@cloudshell:~ (cs571-demo-project-302019) $
```

## 2. Create image and deploy spark to Kubernetes

- Install the NFS Server Provisioner
  - `helm repo add stable https://charts.helm.sh/stable`
  - `helm repo update`

```
singh19566@cloudshell:~ (cs571-demo-project-302019) $ helm repo add stable https://charts.helm.sh/stable
"stable" has been added to your repositories
singh19566@cloudshell:~ (cs571-demo-project-302019) $ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "stable" chart repository
Update Complete. ✨Happy Helming!✨
```

# Implementation



kubernetes

- helm install nfs stable/nfs-server-provisioner\
- set persistence.enabled=true,persistence.size=5Gi

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ helm install nfs stable/nfs-server-provisioner --set persistence.enabled=true,persistence.size=5Gi
WARNING: This chart is deprecated
NAME: nfs
LAST DEPLOYED: Wed Apr 21 18:36:01 2021
NAMESPACE: default
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The NFS Provisioner service has now been installed.

A storage class named 'nfs' has now been created
and is available to provision dynamic volumes.

You can use this storageclass by creating a `PersistentVolumeClaim` with the
correct storageClassName attribute. For example:

---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: test-dynamic-volume-claim
spec:
  storageClassName: "nfs"
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Mi
```

# Implementation



kubernetes

3. Create a persistent disk volume and a pod to use NFS spark-pvc.yaml:

- `kubectl apply -f spark-pvc.yaml`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubectl apply -f spark-pvc.yaml
persistentvolumeclaim/spark-data-pvc created
pod/spark-data-pod created
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

4. Create and prepare your application JAR file

- `docker run -v /tmp:/tmp -it bitnami/spark -- find /opt/bitnami/spark/examples/jars/ -name spark-examples* -exec cp {} /tmp/my.jar \;`

After running the above command, you should see this

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ docker run -v /tmp:/tmp -it bitnami/spark -- find /opt/bitnami/spark/examples/jars/ -name spark-examples* -exec cp {} /tmp/my.jar \;
18:49:01.59
18:49:01.59 Welcome to the Bitnami spark container
18:49:01.59 Subscribe to project updates by watching https://github.com/bitnami/bitnami-docker-spark
18:49:01.59 Submit issues and feature requests at https://github.com/bitnami/bitnami-docker-spark/issues
18:49:01.59
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

# Implementation



kubernetes

5. Add a test file with a line of words that we will be using later for the word count test
  - `echo "how much wood could a woodpecker chuck if a woodpecker could chuck wood">/tmp/test.txt`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ echo "how much wood could a woodpecker chuck if a woodpecker could chuck wood" > /tmp/test.txt
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

6. Copy the JAR file containing the application, and any other required files, to the PVC using the mount point
  - `kubect cp /tmp/my.jar spark-data-pod:/data/my.jar`
  - `kubect cp /tmp/test.txt spark-data-pod:/data/test.txt`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubect cp /tmp/my.jar spark-data-pod:/data/my.jar
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubect cp /tmp/test.txt spark-data-pod:/data/test.txt
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

7. Make sure the files are inside the persistent volume
  - `kubect exec -it spark-data-pod -- ls -al /data`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubect exec -it spark-data-pod -- ls -al /data
total 1504
drwxrwsrwx 2 root root    4096 Apr 21 18:53 .
drwxr-xr-x 1 root root    4096 Apr 21 18:45 ..
-rw-r--r-- 1 1001 root 1527168 Apr 21 18:53 my.jar
-rw-r--r-- 1 1000 1001     72 Apr 21 18:53 test.txt
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

# Implementation



kubernetes

8. Deploy Apache Spark on Kubernetes using the shared volume spark-chart. yaml:

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ vi spark-chart.yaml
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ cat spark-chart.yaml
service:
  type: LoadBalancer
worker:
  replicaCount: 3
  extraVolumes:
    - name: spark-data
      persistentVolumeClaim:
        claimName: spark-data-pvc
  extraVolumeMounts:
    - name: spark-data
      mountPath: /data
singhl9566@cloudshell:~ (cs571-demo-project-302019)$
```

9. Check the pods is running:

- [kubectl get pods](#)

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nfs-nfs-server-provisioner-0        1/1     Running   0           23m
spark-data-pod                      1/1     Running   0           13m
singhl9566@cloudshell:~ (cs571-demo-project-302019)$
```

10. Deploy Apache Spark on the Kubernetes cluster using the Bitnami Apache Spark Helm chart and supply it with the configuration file above

- [helm repo add bitnami https://charts.bitnami.com/bitnami](#)  
- [helm install spark bitnami/spark -f spark-chart.yaml](#)

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ helm repo add bitnami https://charts.bitnami.com/bitnami
"bitnami" has been added to your repositories
singhl9566@cloudshell:~ (cs571-demo-project-302019)$
```



# Implementation



kubernetes

11. Get the external IP of the running pod
  - `kubectl get svc -l "app.kubernetes.io/instance=spark,app.kubernetes.io/name=spark"`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubectl get svc -l "app.kubernetes.io/instance=spark,app.kubernetes.io/name=spark"
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)                                     AGE
spark-headless      ClusterIP     None         <none>         <none>                                     56s
spark-master-svc    LoadBalancer 10.3.244.35   35.247.40.216 7077:32060/TCP,80:30975/TCP               56s
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```

12. Open the external ip on your browser,

Spark Master at spark://spark-master-0.spark-headless.default.svc.cluster.local:7077

URL: spark://spark-master-0.spark-headless.default.svc.cluster.local:7077

Alive Workers: 1

Cores in use: 1 Total, 0 Used

Memory in use: 14.6 GiB Total, 0.0 B Used

Resources in use:

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker id	Address	State	Cores	Memory	Resources
worker-20210421191926-10.0.0.5-44867	10.0.0.5:44867	ALIVE	1 (0 Used)	14.6 GiB (0.0 B Used)	

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------



kubernetes

# Table of Content

- ❖ **Introduction**
- ❖ **Design**
  - Project Description
  - Concepts - Basic Ideas
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**



kubernetes

# Test Result - Wordcount on Spark

## → Submit a word count task :

- `kubectrl run --namespace default spark-client --rm --tty -i --restart='Never' \`  
`--image docker.io/bitnami/spark:3.0.1-debian-10-r115 \`  
`-- spark-submit --master spark://LOAD-BALANCER-External-ip- ADDRESS:7077 \ --deploy-mode cluster \`  
`--class org.apache.spark.examples.JavaWordCount \`  
`/data/my.jar /data/test.txt`

You should see something like this after the above command

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ kubectrl run --namespace default spark-client --rm --tty -i --restart='Never' \
> --image docker.io/bitnami/spark:3.0.1-debian-10-r115 \
> -- spark-submit --master spark://35.247.40.216:7077 \
> --deploy-mode cluster \
> --class org.apache.spark.examples.JavaWordCount \
> /data/my.jar /data/test.txt
If you don't see a command prompt, try pressing enter.
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.NativeCodeLoader).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
21/04/21 19:39:56 INFO SecurityManager: Changing view acls to: spark
21/04/21 19:39:56 INFO SecurityManager: Changing modify acls to: spark
21/04/21 19:39:56 INFO SecurityManager: Changing view acls groups to:
21/04/21 19:39:56 INFO SecurityManager: Changing modify acls groups to:
21/04/21 19:39:56 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set()
21/04/21 19:39:57 INFO Utils: Successfully started service 'driverClient' on port 43135.
21/04/21 19:39:57 INFO TransportClientFactory: Successfully created connection to /35.247.40.216:7077 after 67 ms (0 ms spent in boot)
21/04/21 19:39:57 INFO ClientEndpoint: Driver successfully submitted as driver-20210421193957-0000
21/04/21 19:39:57 INFO ClientEndpoint: ... waiting before polling master for driver state
21/04/21 19:40:02 INFO ClientEndpoint: ... polling master for driver state
21/04/21 19:40:02 INFO ClientEndpoint: State of driver-20210421193957-0000 is RUNNING
21/04/21 19:40:02 INFO ClientEndpoint: Driver running on 10.0.2.4:35165 (worker-20210421192047-10.0.2.4-35165)
21/04/21 19:40:02 INFO ShutdownHookManager: Shutdown hook called
21/04/21 19:40:02 INFO ShutdownHookManager: Deleting directory /tmp/spark-52b94eeb-dd8f-432d-acda-33fa26df2b9d
pod "spark-client" deleted
singhl9566@cloudshell:~ (cs571-demo-project-302019)$
```

# Test Result - Wordcount on Spark



kubernetes

→ And on your browser, you should see this task finished

← → ↻ ⚠ Not Secure | 35.247.40.216 ☆ 🔒 🔍 🏠 👤 ⋮

## Spark Master at spark://spark-master-0.spark-headless.default.svc.cluster.local:7077

URL: spark://spark-master-0.spark-headless.default.svc.cluster.local:7077

Alive Workers: 3

Cores in use: 3 Total, 0 Used

Memory in use: 43.9 GiB Total, 0.0 B Used

Resources in use:

Applications: 0 Running, 1 Completed

Drivers: 0 Running, 1 Completed

Status: ALIVE

### Workers (3)

Worker Id	Address	State	Cores	Memory	Resources
<a href="#">worker-20210421191926-10.0.0.5-44867</a>	10.0.0.5:44867	ALIVE	1 (0 Used)	14.6 GiB (0.0 B Used)	
<a href="#">worker-20210421192047-10.0.2.4-35165</a>	10.0.2.4:35165	ALIVE	1 (0 Used)	14.6 GiB (0.0 B Used)	
<a href="#">worker-20210421192120-10.0.1.9-33157</a>	10.0.1.9:33157	ALIVE	1 (0 Used)	14.6 GiB (0.0 B Used)	

### Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

### Running Drivers (0)

Submission ID	Submitted Time	Worker	State	Cores	Memory	Resources	Main Class	Duration
---------------	----------------	--------	-------	-------	--------	-----------	------------	----------

### Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
<a href="#">app-20210421194002-0000</a>	JavaWordCount	2	1024.0 MiB		2021/04/21 19:40:02	spark	FINISHED	13 s

### Completed Drivers (1)

Submission ID	Submitted Time	Worker	State	Cores	Memory	Resources	Main Class
<a href="#">driver-20210421193957-0000</a>	2021/04/21 19:39:57	<a href="#">worker-20210421192047-10.0.2.4-35165</a>	FINISHED	1	1024.0 MiB		org.apache.spark.examples.JavaWordCount

# Test Result - Wordcount on Spark



## → View the output of the completed jobs

- ◆ On the browser, you should see the worker node ip address of the finished task

### → Completed Drivers (1)

Submission ID	Submitted Time	Worker	State	Cores	Memory	Resources	Main Class
driver-20210421193957-0000	2021/04/21 19:39:57	<a href="#">worker-20210421192047-10.0.2.4-35165</a>	FINISHED	1	1024.0 MiB		org.apache.spark.examples.JavaWordCount

## → For example, my worker node ip address is 10.0.2.4

- ◆ Get the name of the worker node
  - `kubectl get pods -o wide | grep WORKER-NODE-ADDRESS`
  - `kubectl get pods -o wide | grep 10.0.2.4`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubectl get pods -o wide | grep 10.0.2.4
spark-worker-1          1/1      Running   0           24m    10.0.2.4    gke-spark-default-pool-f4a792fd-32rx    <none>
singh19566@cloudshell:~ (cs571-demo-project-302019)$
```



kubernetes

# Test Result - Wordcount on Spark

→ Execute this pod and see the result of the finished tasks

- `kubectrl exec -it <worker node name> -- bash`
- `kubectrl exec -it spark-worker-1 -- bash`

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ kubectrl exec -it spark-worker-1 -- bash
I have no name!@spark-worker-1:/opt/bitnami/spark$ cd /opt/bitnami/spark/work
```

- `cd /opt/bitnami/spark/work`
- `cat <taskname>/stdout`

```
singhl9566@cloudshell:~ (cs571-demo-project-302019)$ kubectrl exec -it spark-worker-1 -- bash
I have no name!@spark-worker-1:/opt/bitnami/spark$ cd /opt/bitnami/spark/work
I have no name!@spark-worker-1:/opt/bitnami/spark/work$ cat driver-20210421193957-0000 /stdout
cat: driver-20210421193957-0000: Is a directory
cat: /stdout: No such file or directory
I have no name!@spark-worker-1:/opt/bitnami/spark/work$ cat driver-20210421193957-0000/stdout
if: 1
a: 2
how: 1
could: 2
wood: 2
woodpecker: 2
much: 1
chuck: 2
I have no name!@spark-worker-1:/opt/bitnami/spark/work$
```



kubernetes

# Test Result - PageRank on PySpark on the pods

- Execute the spark master pods
  - `kubectl exec -it spark-master-0 -- bash`
- Start pyspark
  - `pyspark`

```
singh19566@cloudshell:~ (cs571-demo-project-302019)$ kubectl exec -it spark-master-0 -- bash
I have no name!@spark-master-0:/opt/bitnami/spark$ pyspark
Python 3.6.13 (default, Apr 19 2021, 18:12:00)
[GCC 8.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
21/04/21 19:50:51 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to

  ____      _
 / ___|  __| | | |
 \___ \  | | | | | |
  ___) | | | | | | |
 |____|_|_|_|_|_|_|_|

 version 3.1.1

Using Python version 3.6.13 (default, Apr 19 2021 18:12:00)
Spark context Web UI available at http://spark-master-0.spark-headless.default.svc.cluster.local:4040
Spark context available as 'sc' (master = local[*], app id = local-1619034653797).
SparkSession available as 'spark'.
>>> █
```



# Test Result - PageRank on PySpark on the pods

- ➔ **Exit pyspark with**
  - `exit()`
- ➔ **Go to the directory where pagerank.py located**
  - `cd /opt/bitnami/spark/examples/src/main/python`
- ➔ **Run the page rank using pyspark**
  - `spark-submit pagerank.py /opt 2`

Note, /opt is an example directory and 2 is the number of iterations you want the page rank to run, we can also change to any numbers, here is my output of running the page rank for directory /opt with 2 iterations

```
CLOUD SHELL
Terminal (cs571-demo-project-302019) x +

file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/greengrass/2017-06-07
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/boto3-1.17.53.dist-info
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/awsccli/examples/ec2
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/pandas/tests/indexes/period
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/personalize-events/2018-03-22
file:/opt/bitnami/python/lib/python3.6/test/cjkencodings
file:/opt/bitnami/java/demo/applets/wireframe
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/cur/2017-01-06
file:/opt/bitnami/spark/examples/src/main/java/org/apache/spark/examples/mllib
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/appstream/2016-12-01
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/secretsmanager/2017-10-17
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/awsccli/examples/redshift
file:/opt/bitnami/spark/python/pyspark/python/pyspark
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/elb/2012-06-01
file:/opt/bitnami/spark/venv/lib/python3.6/site-packages/botocore/data/ec2/2015-03-01

If provided paths are partition directories, please set "basePath" in the options of the data source to specify the root
at scala.Predef$.assert(Predef.scala:223)
at org.apache.spark.sql.execution.datasources.PartitioningUtils$.parsePartitions(PartitioningUtils.scala:172)
at org.apache.spark.sql.execution.datasources.PartitioningUtils$.parsePartitions(PartitioningUtils.scala:104)
at org.apache.spark.sql.execution.datasources.PartitioningAwareFileIndex.inferPartitioning(PartitioningAwareFileIndex.scala:73)
at org.apache.spark.sql.execution.datasources.PartitioningAwareFileIndex.inferPartitioning(PartitioningAwareFileIndex.scala:167)
at org.apache.spark.sql.execution.datasources.DataSource.resolveRelation(DataSource.scala:418)
at org.apache.spark.sql.DataFrameReader.loadV1Source(DataFrameReader.scala:326)
at org.apache.spark.sql.DataFrameReader.load(DataFrameReader.scala:308)
at scala.Option.getOrElse(Option.scala:189)
at org.apache.spark.sql.DataFrameReader.load(DataFrameReader.scala:308)
at org.apache.spark.sql.DataFrameReader.text(DataFrameReader.scala:945)
at sun.reflect.NativeMethodAccessorImpl.invoke(Native Method)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:62)
at java.lang.reflect.Method.invoke(Method.java:498)
at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
at py4j.Gateway.invoke(Gateway.java:282)
at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
at py4j.commands.CallCommand.execute(CallCommand.java:79)
at py4j.GatewayConnection.run(GatewayConnection.java:238)
at java.lang.Thread.run(Thread.java:748)

21/04/21 19:54:29 INFO SparkContext: Invoking stop() from shutdown hook
21/04/21 19:54:29 INFO SparkUI: Stopped Spark web UI at http://spark-master-0:spark-headless.default.svc.cluster.local:
21/04/21 19:54:29 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
21/04/21 19:54:29 INFO MemoryStore: MemoryStore cleared
21/04/21 19:54:29 INFO BlockManager: BlockManager stopped
21/04/21 19:54:29 INFO BlockManagerMaster: BlockManagerMaster stopped
21/04/21 19:54:29 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
21/04/21 19:54:29 INFO SparkContext: Successfully stopped SparkContext
21/04/21 19:54:29 INFO ShutdownHookManager: Shutdown hook called
21/04/21 19:54:29 INFO ShutdownHookManager: Deleting directory /tmp/spark-f4ed5100-20cb-4be4-ae31-07e3d5e7390e/pyspark-c
21/04/21 19:54:29 INFO ShutdownHookManager: Deleting directory /tmp/spark-f4ed5100-20cb-4be4-ae31-07e3d5e7390e
21/04/21 19:54:29 INFO ShutdownHookManager: Deleting directory /tmp/spark-370e4c27-07aa-40dc-af44-65dcbce2af5b
I have no name!@spark-master-0:/opt/bitnami/spark/examples/src/main/python$
```



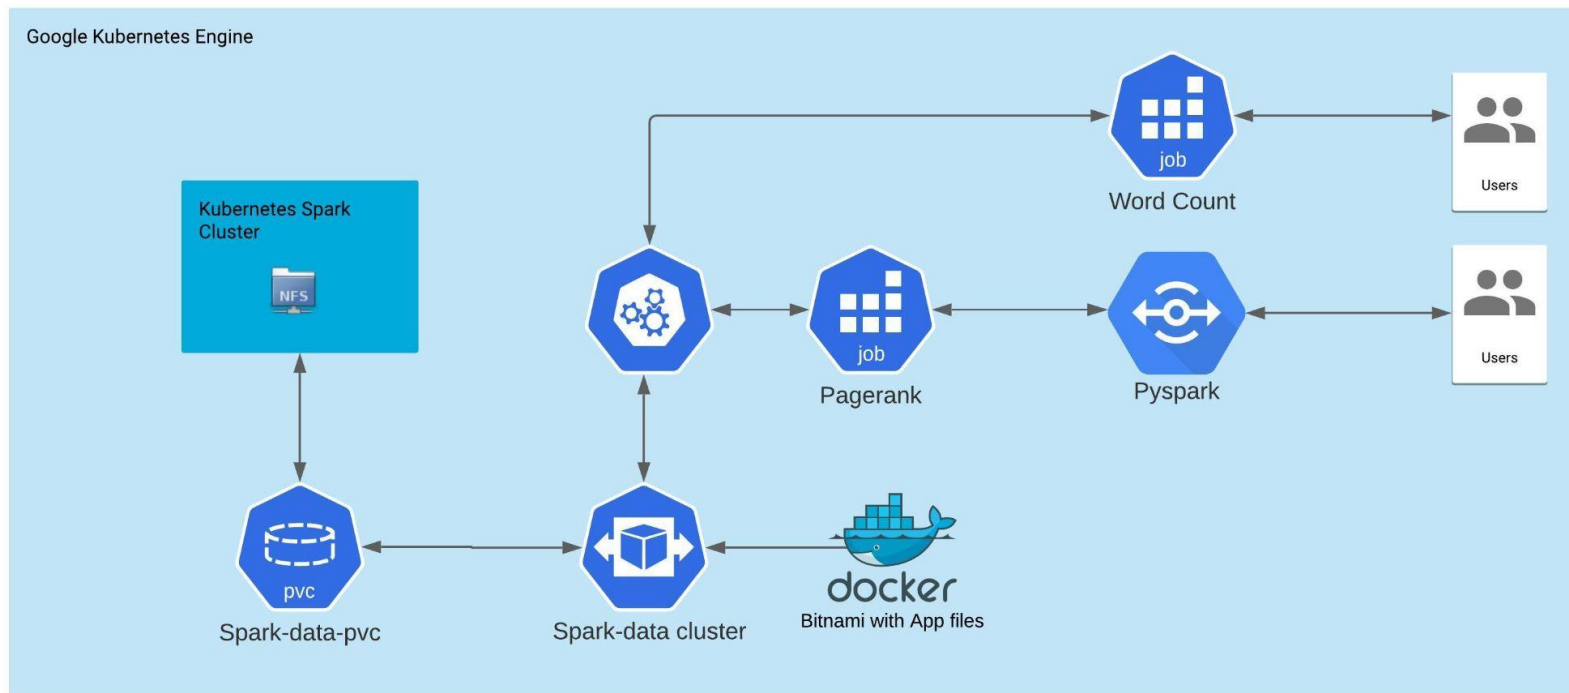
# Table of Content

- ❖ **Introduction**
- ❖ **Design**
  - Project Description
  - Concepts - Basic Ideas
    - Workcount
    - PageRank
    - PageRank Calculation Example
- ❖ **Implementation**
- ❖ **Test Results**
  - Wordcount
  - PageRank
- ❖ **Schematic of the Project**
- ❖ **Conclusion**
- ❖ **Bibliography/References**

# Schematic of Project



kubernetes



# Conclusion

- ★ Apache Spark is a cluster computing platform designed to be fast, speed side and extends the popular MapReduce model to efficiently supports more type of computations, including interactive queries and stream processing.
- ★ Since Spark integrates closely with other big data tool, hence this tight integration is the ability to build an application that seamlessly combines different computation model.
- ★ Spark is also highly fault-tolerant; if one node fails, the failed tasks are distributed across the other nodes.
- ★ **The entire project is developed in Google Cloud Platform which provides different tools from which we are using Google Kubernetes Engine(GKE).**
- ★ **Then using PySpark API we implemented Word Count and PageRank applications on Apache Spark running on Google Kubernetes Engine.**

# Bibliography/References

- [https://npu85.npu.edu/~henry/npu/classes/mapreduce/word\\_count/slide/wordcount\\_explain.html](https://npu85.npu.edu/~henry/npu/classes/mapreduce/word_count/slide/wordcount_explain.html)
- <https://kubernetes.io/blog/2018/03/apache-spark-23-with-native-kubernetes/>
- <https://databricks.com/blog/2018/03/06/apache-spark-2-3-with-native-kubernetes-support.html>
- <https://spark.apache.org/docs/latest/api/python/index.html>
- [https://npu85.npu.edu/~henry/npu/classes/learning\\_spark/key\\_value\\_pair/slide/Example\\_PageRank.html](https://npu85.npu.edu/~henry/npu/classes/learning_spark/key_value_pair/slide/Example_PageRank.html)
- <https://www.sciencedirect.com/topics/computer-science/apache-spark>
- [https://stanford.edu/~rezab/sparkclass/slides/itas\\_workshop.pdf](https://stanford.edu/~rezab/sparkclass/slides/itas_workshop.pdf)